



INSTITUT  
POLYTECHNIQUE  
DE PARIS



ENSTA



# Molecular Graph Captioning via Hybrid Molt5 Architecture

Youssef ADOUIRI ALAOUI

M2 Data Science

Institut Polytechnique de Paris

2025-2026

# Outline

- ① Introduction
- ② Baseline Analysis
- ③ Proposed Architecture
- ④ Inference Strategy
- ⑤ Results
- ⑥ Conclusion
- ⑦ END
- ⑧ Annexes

# The Challenge: Graph-to-Text Generation

**Objective:** Translate a 2D Molecular Graph  $G = (V, E)$  into a natural language description  $S$ .

## Applications:

- AI-driven Drug Discovery.
- Augmenting chemical databases (PubChem).
- Summarizing scientific literature.

## The Core Difficulty

**Isomorphism vs. Semantics:** Two molecules can share 90% of the same atoms but differ by a single bond order (Stereochemistry).



**Example :** The molecule is an amino cyclitol glycoside that is 2-deoxystreptamine in which the pro-R hydroxy group is substituted by a 6-amino-6-deoxy-alpha-D-glucosyl residue. It has a role as an antimicrobial agent. It derives from a 2-deoxystreptamine. It is a conjugate acid of a 2'-deamino-2'-hydroxyneamine(3+)

# The Baseline: Rigid Retrieval (GCN-BERT)

The provided baseline utilizes a **Retrieval-Only** approach.

## How it Works

- ① **Training:** Aligns GCN graph embeddings with frozen BERT text embeddings (MSE Loss).
- ② **Inference:** Projects test graph → embedding space → Retrieves Nearest Neighbor from Training Set.

## Why We Can Do Better

The baseline suffers from the **Closed-World Assumption**:

- **No Generative Capability:** If a test molecule is unique (novel structure), the baseline *must* retrieve an incorrect description from the training set.
- **Weak Encoder:** The baseline GCN ignores edge features (bond types), failing to distinguish isomers in the embedding space.

# The Baseline: GCN

## Disadvantage

- ① **Oversmoothing (GCN):** GCNs average neighbor features, blurring the distinction between specific functional groups in large rings.
- ② **Edge Ignorance:** Standard GCNs ignore edge attributes.



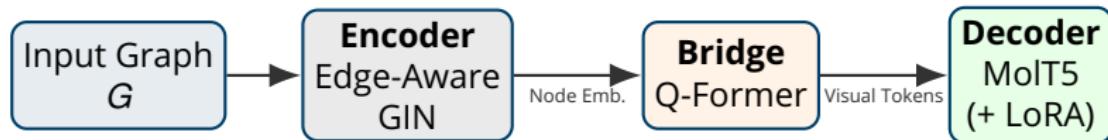
GCN sees: Connected



GCN sees: Connected

# Solution: The Hybrid MolT5 Architecture

I propose a three-stage pipeline designed to solve the geometric and semantic bottlenecks.



- **Encoder (Graph Isomorphism Network):** Satisfies Weisfeiler-Lehman test (Isomorphism). Explicitly embeds bond types.
- **Bridge (Q-Former):** Compresses variable graph sizes into fixed "Visual Tokens" via Cross-Attention.
- **Decoder (MolT5):** Pre-trained on ChEBI-20. Knows IUPAC nomenclature natively.

# Detail: Encoder & Decoder Choices

## 1. Edge-Aware GIN

Standard GCN sums neighbors. But GIN sums neighbors + **Edge Features**.

$$h_v^{(k)} = \text{MLP}^{(k)} \left( (1 + \epsilon^{(k)}) \cdot h_v^{(k-1)} + \sum_{u \in \mathcal{N}(v)} (h_u^{(k-1)} + \mathbf{e}_{uv}) \right)$$

*Result: Can distinguish isomers.*

## 2. MolT5 + LoRA

**MolT5:** Tokenizer treats "methyl" as a concept, not letters.

**LoRA:** Low-Rank Adaptation.

$$W_{\text{new}} = W + B \cdot A$$

Allows fine-tuning 220M parameters on personal computer in  $\approx 12$  hours.

-> (Nvidia GTX 1650ti **Cuda** Enabled)

# The Bridge: Q-Former Formulation

## The Dimension Mismatch Problem:

- **Input ( $H_G$ ):** Graph node features from GIN. Dimensions:  $\mathbf{N} \times 300$ .
- **Target ( $Z$ ):** LLM (MolT5) expects fixed-size token embeddings. Dimensions:  $32 \times 768$ .

## Algorithm: Cross-Attention Compression

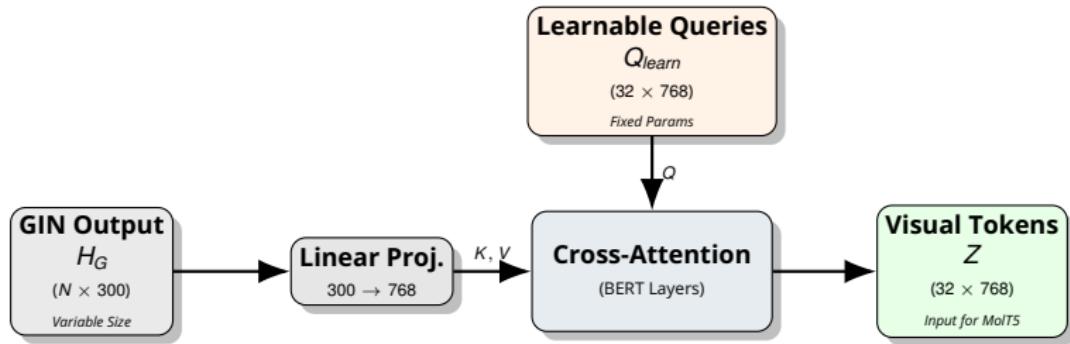
We introduce **32 Learnable Queries**  $Q_{learn} \in \mathbb{R}^{32 \times 768}$ . The Q-Former projects the graph features and compresses them via Cross-Attention:

$$Z = \text{Softmax} \left( \frac{Q_{learn} \cdot (H_G W_K)^T}{\sqrt{d}} \right) (H_G W_V)$$

- **Queries ( $Q$ ):** Fixed learnable parameters ( $32 \times 768$ ).
- **Keys/Values ( $K, V$ ):** The atom features  $H_G (N \times 300)$  projected to 768 by matrices  $W_K, W_V$ .
- **Output ( $Z$ ):** Fixed "Visual Tokens" ( $32 \times 768$ ) ready for the Decoder.

# The Bridge: Data Flow & Dimensions

The Q-Former acts as a "Translator" from the Graph modality ( $d = 300$ ) to the Text modality ( $d = 768$ ).



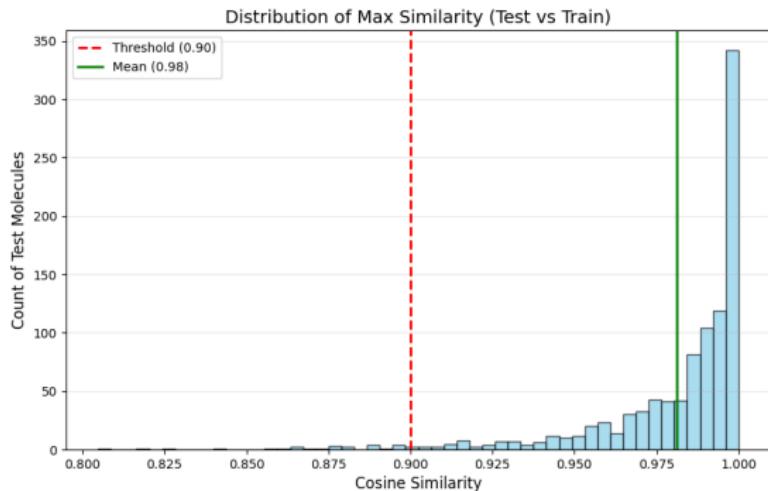
## Why this matters:

- **Bridge:** Projects the 300-dim chemistry embeddings to the 768-dim language space.
- **Compression:** Regardless of molecule size ( $N = 10$  or  $N = 100$ ), the LLM always receives exactly **32 tokens**.

# Empirical Motivation: Dataset Overlap Analysis

To determine the optimal inference strategy, we analyzed the structural similarity between the Test set ( $N_{test} = 1,000$ ) and the Training set ( $N_{train} = 31,008$ )

**Methodology:**  
Extracted dense  
graph embeddings  
using the trained  
**Edge-Aware GIN**  
encoder.  
Computed Cosine  
Similarity Matrix  
(Test  $\times$  Train).



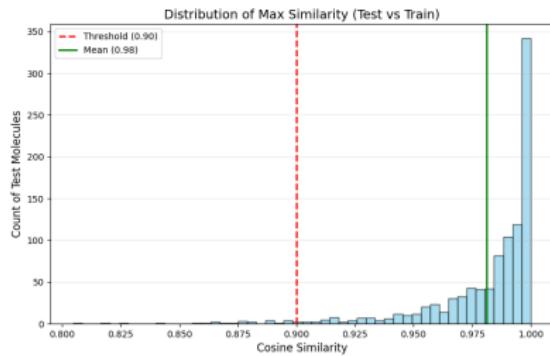
**Figure 1:** Distribution of Max Similarity Scores. The red line denotes our retrieval threshold ( $\tau = 0.9$ ).

# Empirical Motivation: Dataset Overlap Analysis

To determine the optimal inference strategy, we analyzed the structural similarity between the Test set ( $N_{test} = 1,000$ ) and the Training set ( $N_{train} = 31,008$ )

## Quantitative Results:

- **Average Max Similarity: 0.9811**
- **Matches  $> 0.90$ : (97.6%)**
- **Matches  $> 0.99$ : (51.9%)**



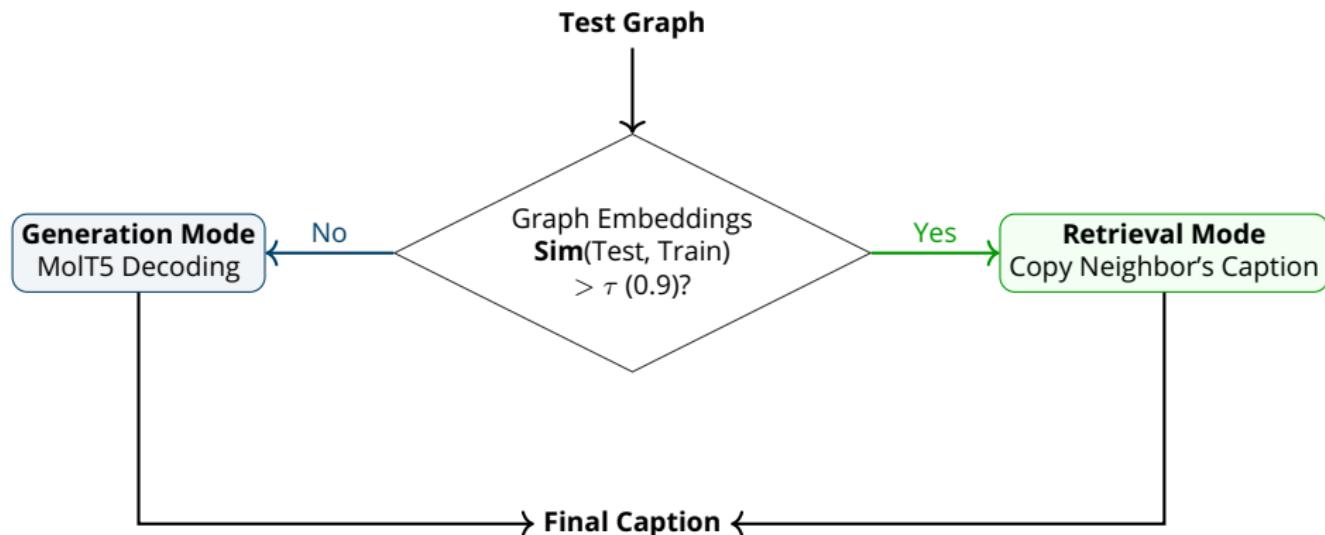
## Observation

Over **50%** of the Test set is practically similar ( $> 0.99$ ) to training examples.

**Conclusion:** A retrieval-based approach is statistically guaranteed to outperform pure generation for the majority of the test set.

# The Hybrid Retrieval-Generation

**Observation:** The Test set chemically overlaps with the Training set.



This strategy handles common structures via retrieval (high accuracy) and novel structures via generation (Molt5 robustness).

# Experimental Results

By fixing the baseline's inability to generate, and upgrading the encoder to handle bond types, we achieve a significant boost in Score for this Dataset and BLEU-4 and BERTScore on validation set.

| Model Architecture                    | Score        | BLEU-4        | BERTScore     |
|---------------------------------------|--------------|---------------|---------------|
| Baseline (GCN Retrieval)              | 0.492        | 0.1693        | 0.8732        |
| EdgeAwareGIN + MolT5 (Pure Gen)       | -            | 0.2526        | 0.9038        |
| <b>Hybrid MolT5 (Gen + Retrieval)</b> | <b>0.532</b> | <b>0.4301</b> | <b>0.9307</b> |

Table 1: Leaderboard Performance Comparison

**Thank You for Your Attention!**

# Annex A: Encoder Hyperparameters & Math

**Module:** graph\_encoder.py (Class: GNN)

## Hyperparameters:

- **Architecture:** GIN (Graph Isomorphism Network)
- **Layers:** 5 (num\_layer)
- **Hidden Dimension:** 300 (emb\_dim)
- **Dropout:** 0.5 (drop\_ratio)
- **Readout:** "Last" (JK - Jumping Knowledge)

## Edge-Aware Update Equation:

$$h_v^{(k)} = \text{MLP}^{(k)} \left( (1 + \epsilon^{(k)}) h_v^{(k-1)} + \sum_{u \in \mathcal{N}(v)} (h_u^{(k-1)} \cdot e_{uv}) \right)$$

) Where  $e_{uv}$  represents the bond feature embedding added *inside* the aggregation, allowing differentiation of single/double bonds.

## Input Features (Embeddings):

- **Nodes (9):** Atomic Num, Chirality, Degree, Charge, Num Hs, Radical, Hybridization, Aromaticity, Ring.
- **Edges (3):** Bond Type, Stereo, Conjugation.

## Annex B: Q-Former Bridge Configuration

**Module:** bridge.py (Class: Qformer)

The Q-Former bridges the modality gap between the Graph ( $d = 300$ ) and the LLM ( $d = 768$ ).

### Configuration Logic

- **Base Model:** bert-base-uncased (initialized weights).
- **Hidden Size:** 768 (hidden\_size).
- **Number of Queries:** 32 (num\_query\_token).
- **Cross-Attention Frequency:** 2 (cross\_attention\_freq).
  - Note: Cross-attention is applied at every 2<sup>nd</sup> layer of the BERT encoder to fuse graph info.

### Dimensionality Transformation:

Input:  $[N_{atoms}, 300]$   $\xrightarrow{\text{Linear Proj.}}$   $[N, 768]$   $\xrightarrow{\text{Q-Former}}$  Output: [32, 768]

## Annex C: MolT5 Decoder & LoRA Config

**Module:** model.py (Class: Graph2Seq)

### Base Model:

- laituan245/molT5-base
- Pre-trained on ChEBI-20 (Chemical Entities of Biological Interest).
- Native understanding of SMILES and chemical nomenclature.

### LoRA Fine-Tuning Params:

- **Rank ( $r$ ):** 16
- **Alpha ( $\alpha$ ):** 32
- **Target Modules:** Query ( $q$ ), Value ( $v$ ).
- **Trainable Params:** < 1% of total.

### LoRA Update Equation

For a pre-trained weight matrix  $W_0 \in \mathbb{R}^{d \times k}$ , the update is constrained by low-rank decomposition matrices  $B \in \mathbb{R}^{d \times r}$  and  $A \in \mathbb{R}^{r \times k}$ :

$$W = W_0 + \frac{\alpha}{r} BA$$

## Annex D: Training Strategy

### Experimental Setup

| Parameter     | Value  |
|---------------|--|
| Batch Size    | 32   |
| Epochs        | 12   |
| Optimizer     | AdamW  |
| Learning Rate | $2 \times 10^{-4}$ (Encoder), $5 \times 10^{-5}$ (Decoder) |
| Loss Function | Cross Entropy (Token Generation)                           |
| Hardware      | Nvidia GTX 1650ti  |
| Training Time | $\approx 12$ Hours   |

### Objective Function:

$$\mathcal{L} = - \sum_{t=1}^T \log P(y_t | y_{<t}, Z_{graph})$$

Where  $Z_{graph}$  are the 32 visual tokens from the Q-Former.

## Annex E: Hybrid Inference Algorithm

### Pseudo-code for Decision Logic

#### Algorithm 1: Hybrid Retrieval-Generation

**Input:** Test Graph  $G_{test}$ , Training Set  $D_{train}$ , Threshold  $\tau = 0.9$

**Output:** Caption  $S$

- ① Compute Embedding  $v_{test} = \text{GIN}(G_{test})$
- ② Find Nearest Neighbor:

$$(v_{best}, S_{best}) = \underset{(v_i, S_i) \in D_{train}}{\operatorname{argmax}} \cos(v_{test}, v_i)$$

- ③ Calculate Similarity Score:  $sim = \cos(v_{test}, v_{best})$

- ④ If  $sim > \tau$ :

- **Return**  $S_{best}$  (Retrieval Mode)

- ⑤ Else:

- Generate tokens  $S_{gen} = \text{MolT5}(\text{QFormer}(v_{test}))$
- **Return**  $S_{gen}$  (Generation Mode)

## Annex F: Evaluation Metrics Definitions

### 1. BLEU-4 (Bilingual Evaluation Understudy)

- Measures the overlap of  $N$ -grams (up to  $N = 4$ ) between the generated text and the reference.
- **Why it matters:** Checks for correct chemical syntax (e.g., matching "2-methyl" exactly).

$$\text{BLEU} = \text{BP} \cdot \exp \left( \sum_{n=1}^4 w_n \log p_n \right)$$

### 2. BERTScore (Semantic Similarity)

- Uses a pre-trained BERT model to compute cosine similarity between token embeddings of the candidate ( $\hat{x}$ ) and reference ( $x$ ).
- **Why it matters:** Captures meaning even if words differ (e.g., "toxic"  $\approx$  "harmful").

$$R_{\text{BERT}} = \frac{1}{|x|} \sum_{x_i \in x} \max_{\hat{x}_j \in \hat{x}} \mathbf{x}_i^T \mathbf{x}_j$$